



MANIPAL UNIVERSITY
JAIPUR

Operating System – AI2202 (AY: 2024–25)

Dr. Shail Saharan
Assistant Professor
Department of AIML, School of Computer Science and Engineering
Manipal University Jaipur.



Week 1 : Introduction to OS

Overview of Operating System.

Outline

- ❖ What Operating Systems Do
 - ❖ Computer-System Organization and Architecture
 - ❖ Operating-System Structure
 - ❖ Operating-System Operations
 - ❖ Process Management
 - ❖ Memory Management
 - ❖ Storage Management
 - ❖ Protection and Security
 - ❖ Kernel Data Structures
 - ❖ Computing Environments
 - ❖ Open-Source Operating Systems
-

Objectives

- To describe the basic organization of computer systems.
- To provide a grand tour of the major components of operating systems.
- To give an overview of the many types of computing environments.
- To explore several open-source operating systems.


What is an Operating System?

- A program that acts as an intermediary between a user of a computer and the computer hardware
- Operating system goals:
 - Execute user programs and make solving user problems easier
 - Make the computer system convenient to use
 - Use the computer hardware in an efficient manner

Computer System Structure

- The computer system can be divided into four components:
 - Hardware – provides basic computing resources
 - CPU, memory, I/O devices
 - Operating system
 - Controls and coordinates the use of the hardware among various applications and users
 - Application programs – define how the system resources are used to solve the computing problems of the users
 - Word processors, compilers, web browsers, database systems, video games
 - Users
 - People, machines, other computers

Four Components of a Computer System



What Operating Systems Do

- Depends on the point of view
 - Users want convenience, **ease of use**, and **good performance**
 - Don't care about **resource utilization**
 - But shared computers such as **mainframe** or **minicomputers** must keep all users happy
 - Users of dedicated systems such as **workstations** have dedicated resources but frequently use shared resources from **servers**
 - Handheld computers are resource-poor, optimized for usability and battery life
 - Some computers have little or no user interface, such as embedded computers in devices and automobiles.
-

Operating System Definition

- OS is a **resource allocator**
 - Manages all resources
 - Decides between conflicting requests for efficient and fair resource use
- OS is a **control program**
 - Controls execution of programs to prevent errors and improper use of the computer

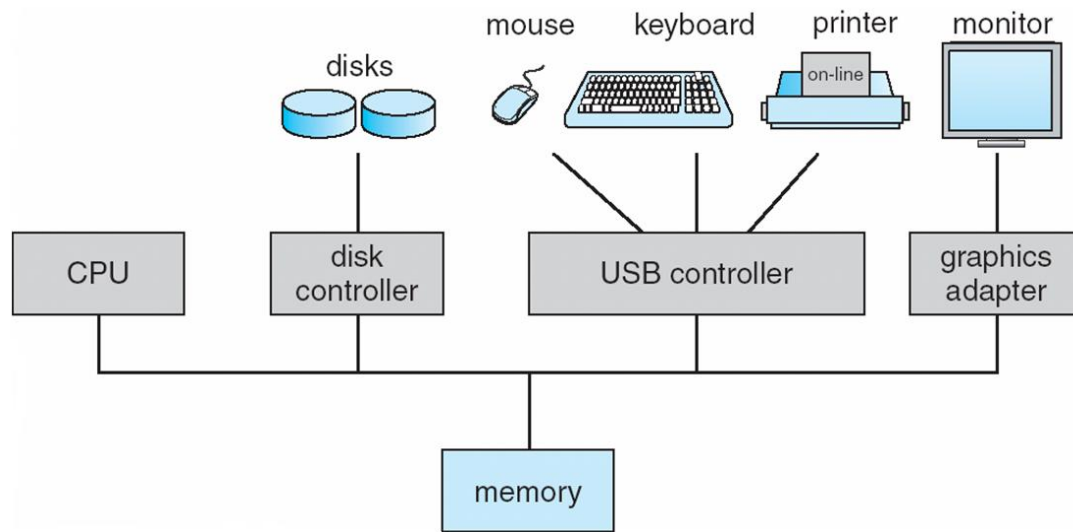
Operating System Definition (Cont.)

- No universally accepted definition
- “Everything a vendor ships when you order an operating system” is a good approximation
 - But varies wildly
- “The one program running at all times on the computer” is the **kernel**.
- Everything else is either
 - a system program (ships with the operating system), or
 - an application program.

Computer Startup

- **Bootstrap program** is loaded at power-up or reboot
 - Typically stored in ROM or EPROM, generally known as **firmware**
 - Initializes all aspects of the system
 - Loads operating system kernel and starts execution

Computer System Organization



- Computer-system operation
 - One or more CPUs, and device controllers connect through a common bus providing access to shared memory
 - Concurrent execution of CPUs and devices competing for memory cycles

Computer-System Operation

- I/O devices and the CPU can execute concurrently.
 - Each device controller oversees a particular device type.
 - Each device controller has a local buffer.
 - CPU moves data from/to main memory to/from local buffers.
 - I/O is the device to the local buffer of the controller.
 - Device controller informs CPU that it has finished its operation by causing an **interrupt**.
-

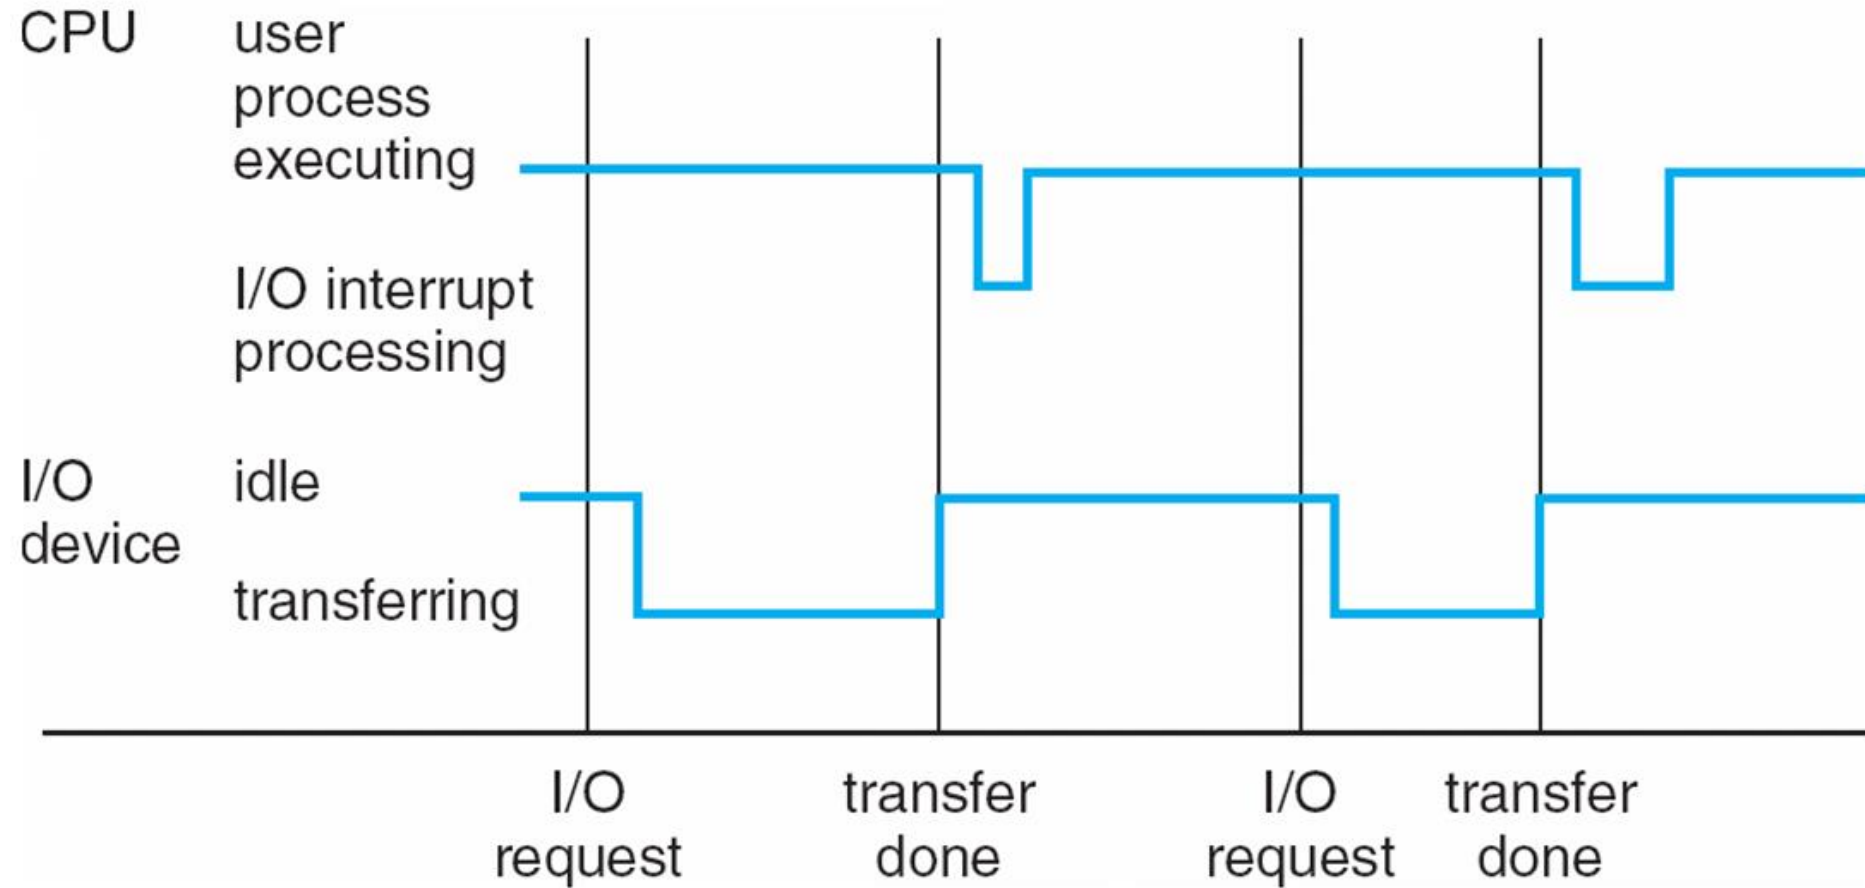
Common Functions of Interrupts

- Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines
- Interrupt architecture must save the address of the interrupted instruction
- A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request
- An operating system is **interrupt-driven**

Interrupt Handling

- The operating system preserves the state of the CPU by storing registers and the program counter
- Determines which type of interrupt has occurred:
 - **polling**
 - **vectored** interrupt system
- Separate segments of code determine what action should be taken for each type of interrupt

Interrupt Timeline



I/O Structure

- After I/O starts, control returns to the user program only upon I/O completion
 - Wait instruction idles the CPU until the next interrupt
 - Wait loop (contention for memory access)
 - At most one I/O request is outstanding at a time, no simultaneous I/O processing

I/O Structure

- After I/O starts, control returns to the user program without waiting for I/O completion
 - **System call** – request to the OS to allow the user to wait for I/O completion
 - **Device-status table** contains an entry for each I/O device indicating its type, address, and state
 - OS indexes into I/O device table to determine device status and to modify table entry to include interrupt

Storage Definitions and Notation Review

- Computer storage, along with most computer throughput, is generally measured and manipulated in bytes and collections of bytes.
 - A kilobyte, or KB, is 1,024 bytes
 - a megabyte, or MB, is 1,024² bytes
 - a gigabyte, or GB, is 1,024³ bytes
 - a terabyte, or TB, is 1,024⁴ bytes
 - a petabyte, or PB, is 1,024⁵ bytes
 - Computer manufacturers often round off these numbers and say that a megabyte is 1 million bytes and a gigabyte is 1 billion bytes.
 - Networking measurements are an exception to this general rule; they are given in bits (because networks move data a bit at a time).
-

Storage Structure

- Main memory – only large storage media that the CPU can access directly
 - Random access
 - Typically volatile
- Secondary storage – an extension of main memory that provides large nonvolatile storage capacity


Storage Structure

- Hard disks – rigid metal or glass platters covered with magnetic recording material
 - Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
 - The **disk controller** determines the logical interaction between the device and the computer
 - **Solid-state disks** – faster than hard disks, nonvolatile
 - Various technologies
 - Becoming more popular
-

Storage Hierarchy

- Storage systems organized in a hierarchy
 - Speed
 - Cost
 - Volatility
- Caching – copying information into a faster storage system; main memory can be viewed as a cache for secondary storage
- Device Driver for each device controller to manage I/O
- Provides a uniform interface between the controller and kernel

Storage-Device Hierarchy



Caching

- Important principle, performed at many levels in a computer (in hardware, operating system, software)
 - Information in use copied from slower to faster storage temporarily
 - Faster storage (cache) checked first to determine if the information is there
 - If it is, information used directly from the cache (fast)
 - If not, data is copied to the cache and used there
 - Cache smaller than storage being cached
 - Cache management important design problem
 - Cache size and replacement policy
-

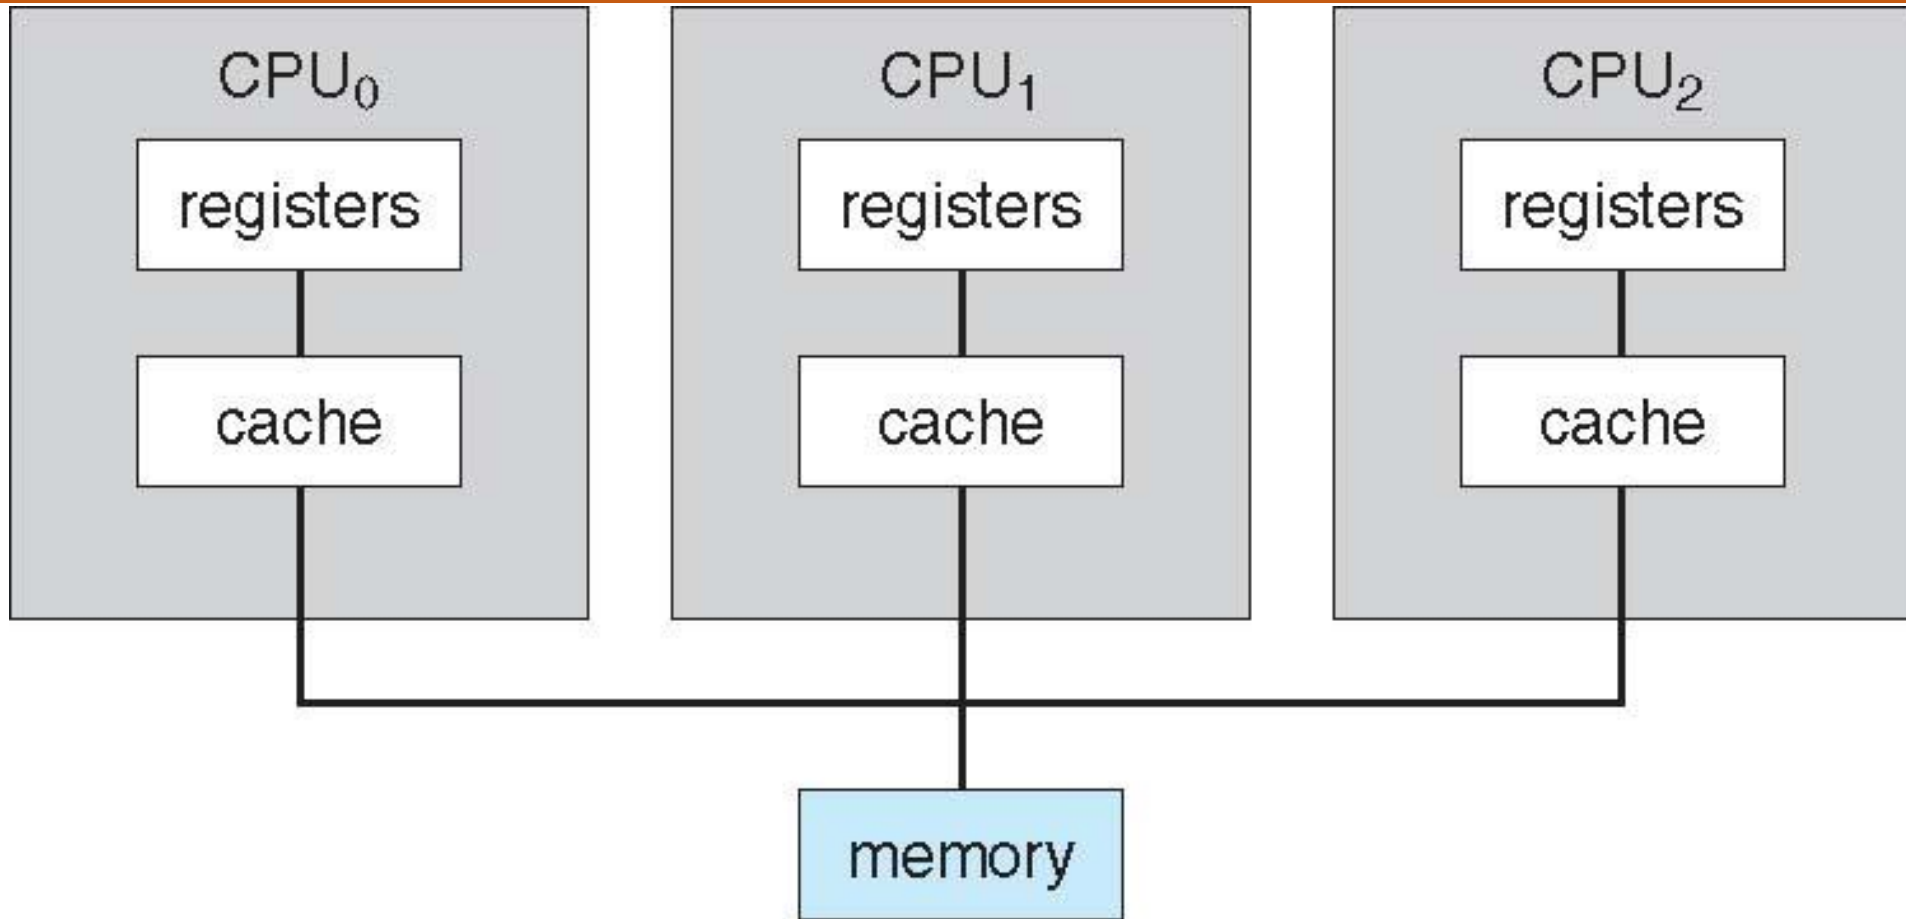
Direct Memory Access Structure

- Used for high-speed I/O devices able to transmit information at close to memory speeds
- Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention
- Only one interrupt is generated per block, rather than one interrupt per byte

Computer-System Architecture


- Most systems use a single general-purpose processor
 - Most systems have special-purpose processors as well
 - **Multiprocessors** systems growing in use and importance
 - Also known as **parallel systems**, **tightly-coupled systems**
 - Advantages include:
 1. **Increased throughput**
 2. **Economy of scale**
 3. **Increased reliability** – graceful degradation or fault tolerance
 - Two types:
 1. **Asymmetric Multiprocessing** – each processor is assigned a specific task.
 2. **Symmetric Multiprocessing** – each processor performs all tasks
-

Symmetric Multiprocessing Architecture



A Dual-Core Design


- Multi-chip and **multicore**
- Systems containing all chips
 - Chassis containing multiple separate systems



Clustered Systems

- Like multiprocessor systems, but multiple systems working together
 - Usually sharing storage via a **storage-area network (SAN)**
 - Provides a **high-availability** service that survives failures
 - **Asymmetric clustering** has one machine in hot-standby mode
 - **Symmetric clustering** has multiple nodes running applications, monitoring each other
 - Some clusters are for **high-performance computing (HPC)**
 - Applications must be written to use **parallelization**
 - Some have **distributed lock manager (DLM)** to avoid conflicting operations

Clustered Systems




Operating System Structure

- **Multiprogramming (Batch system)** needed for efficiency
 - Single user cannot keep CPU and I/O devices busy at all times
 - Multiprogramming organizes jobs (code and data) so the CPU always has one to execute
 - A subset of total jobs in the system is kept in memory
 - One job selected and run via **job scheduling**
 - When it has to wait (for I/O for example), OS switches to another job

Operating System Structure

- **Timesharing (multitasking)** is a logical extension in which the CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing
 - **Response time** should be < 1 second
 - Each user has at least one program executing in the memory \Rightarrow **process**
 - If several jobs ready to run at the same time \Rightarrow **CPU scheduling**
 - If processes don't fit in memory, **swapping** moves them in and out to run
- **Virtual memory** allows the execution of processes not completely in memory


Memory Layout for Multiprogrammed System



Types of Operating Systems


1. Batch Operating System
2. Multi-Programming System
3. Multi-Processing System
4. Multi-Tasking Operating System
5. Time-Sharing Operating System
6. Distributed Operating System
7. Network Operating System

Batch Operating System




Multi-Programming Operating System

Multiprogramming

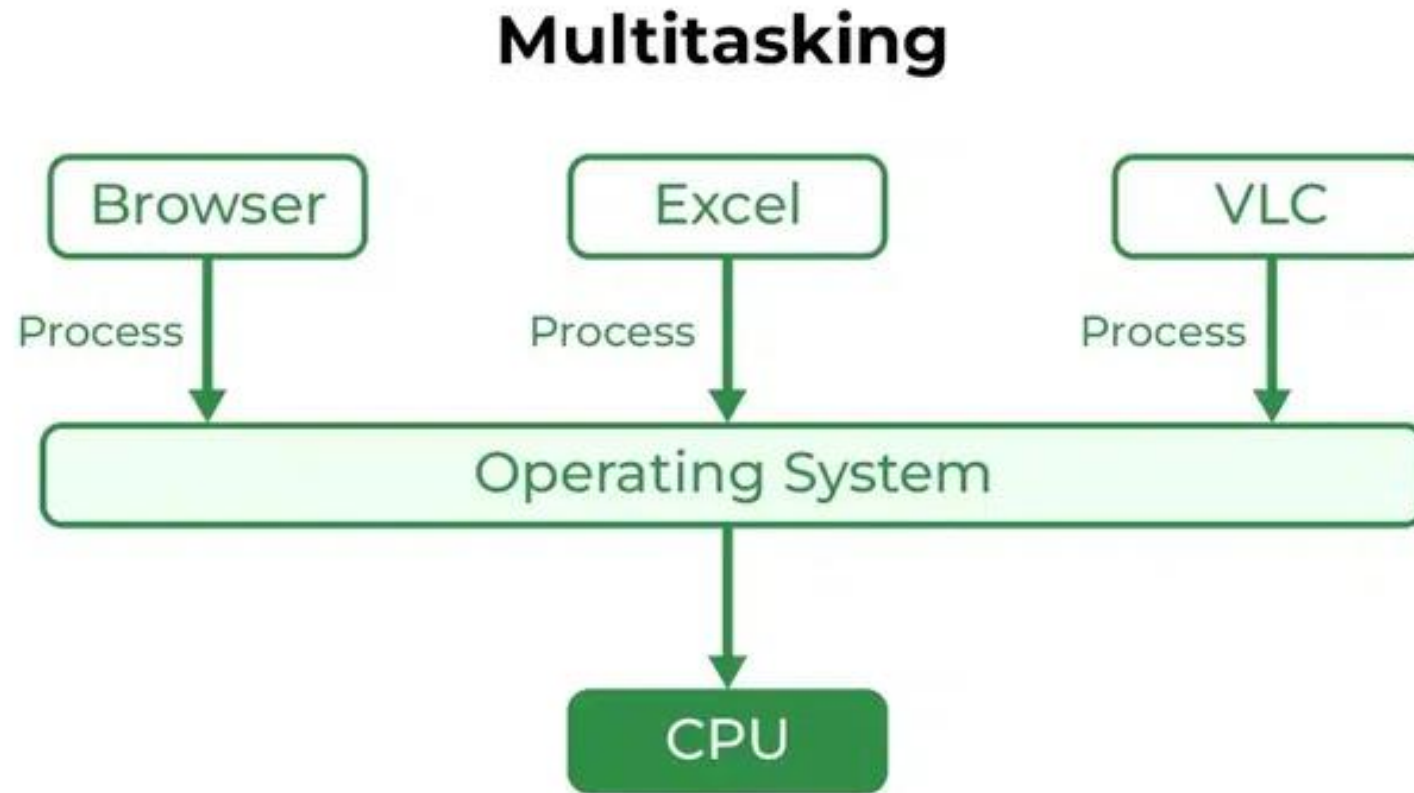


Multi-Processing Operating System

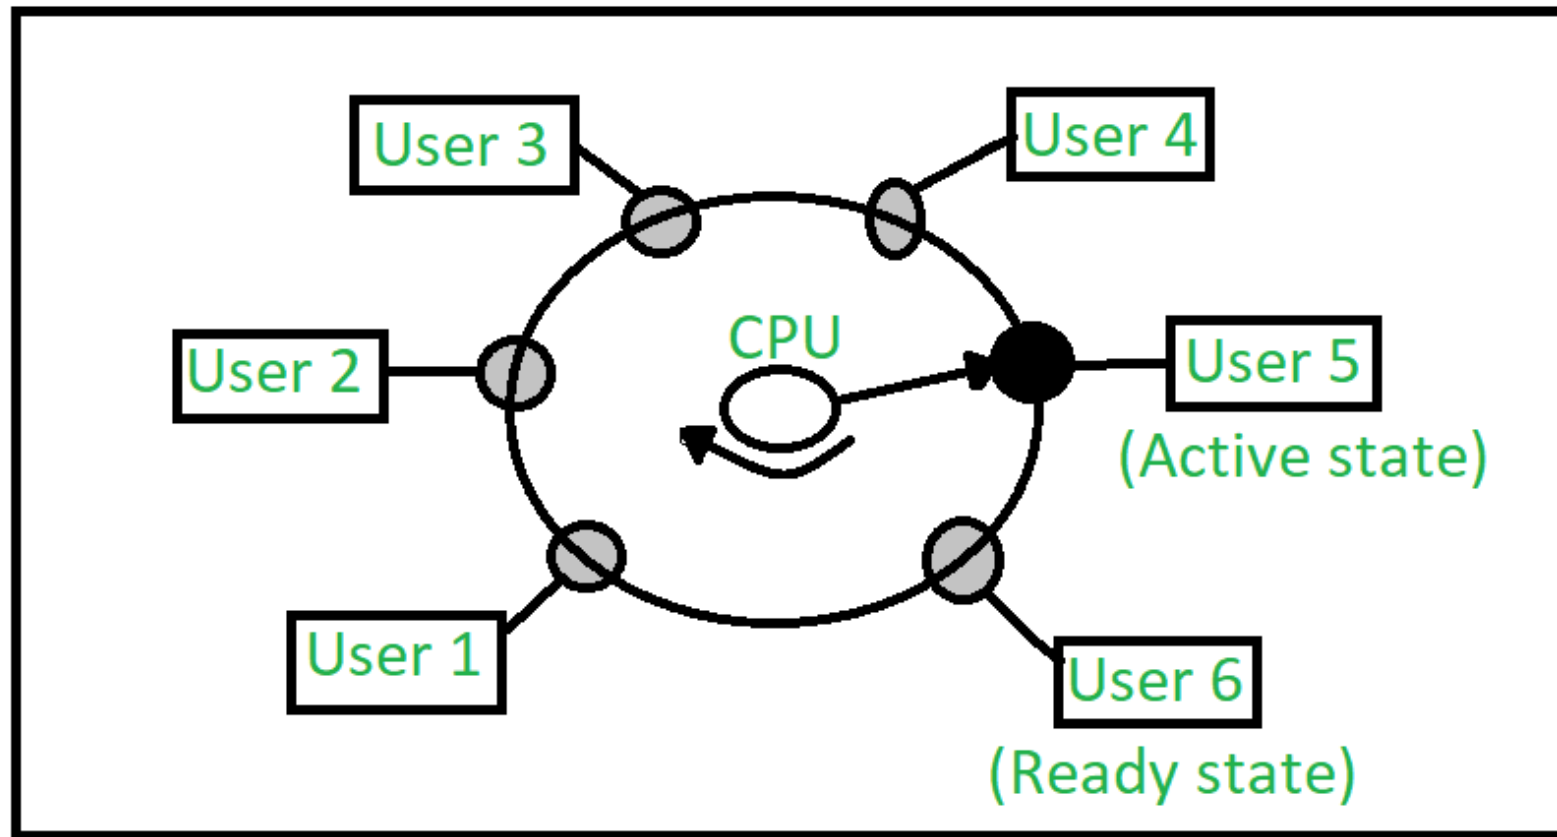
Multiprocessing



Multi-Tasking Operating System

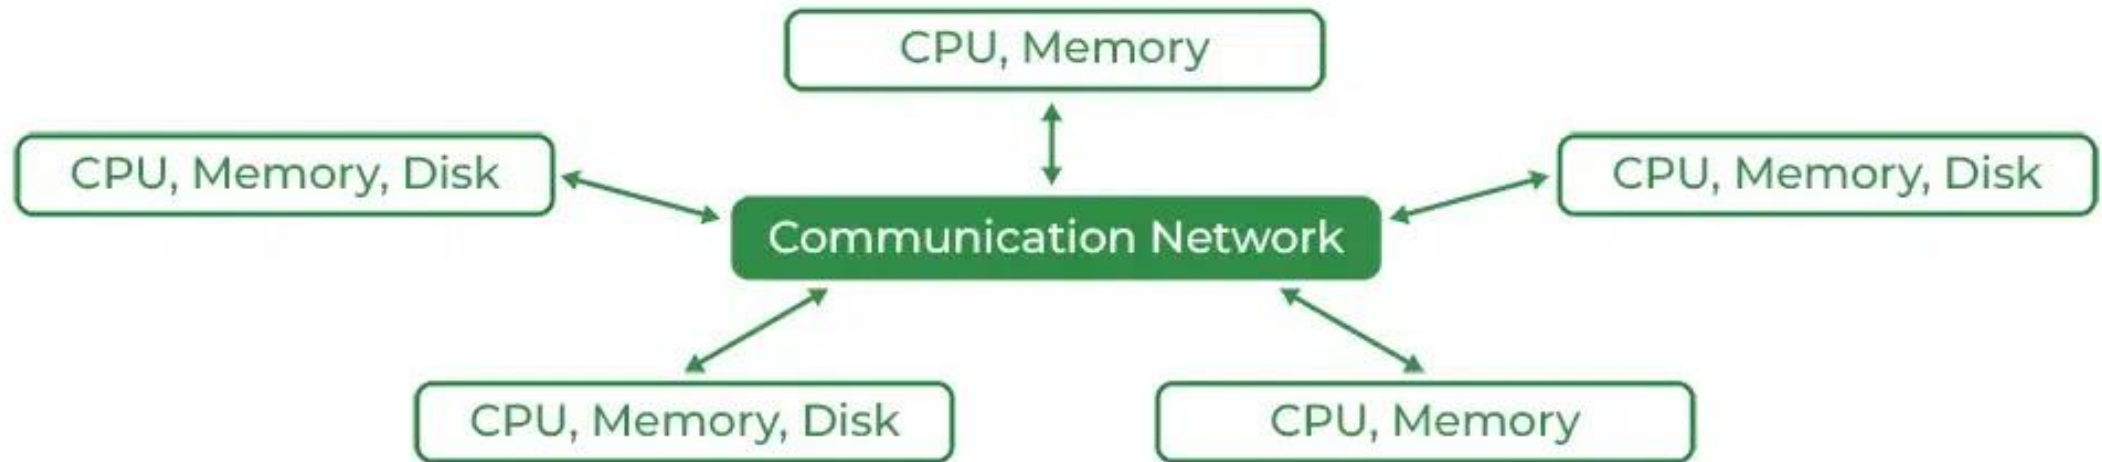


Time-Sharing Operating Systems




Distributed Operating System

Architecture of Distributed OS




Network Operating System




Kernel Data Structures

- Many similar to standard programming data structures


- ***Singly linked list***



- ***Doubly linked list***



- ***Circular linked list***




Kernel Data Structures

- **Binary search tree**


left \leq right

- Search performance is $O(n)$
- **Balanced binary search tree** is $O(\lg n)$



Kernel Data Structures

- **Hash function** can create a **hash map**



- **Bitmap** – string of n binary digits representing the status of n items
- Linux data structures defined in
 include files `<linux/list.h>`, `<linux/kfifo.h>`,
 `<linux/rbtree.h>`

Computing Environments - Traditional

- Stand-alone general-purpose machines
 - But blurred as most systems interconnect with others (i.e., the Internet)
 - **Portals** provide web access to internal systems
 - **Network computers (thin clients)** are like Web terminals
 - Mobile computers interconnect via **wireless networks**
 - Networking becoming ubiquitous – even home systems use **firewalls** to protect home computers from Internet attacks
-

Computing Environments - Mobile


- Handheld smartphones, tablets, etc
- What is the functional difference between them and a “traditional” laptop?
- Extra feature – more OS features (GPS, gyroscope)
- Allows new types of apps like ***augmented reality***
- Use IEEE 802.11 wireless, or cellular data networks for connectivity
- Leaders are **Apple iOS** and **Google Android**

Computing Environments – Distributed

- Distributed computing
 - Collection of separate, possibly heterogeneous, systems networked together
 - **Network** is a communications path, **TCP/IP** most common
 - **Local Area Network (LAN)**
 - **Wide Area Network (WAN)**
 - **Metropolitan Area Network (MAN)**
 - **Personal Area Network (PAN)**
 - **Network Operating System** provides features between systems across the network
 - Communication scheme allows systems to exchange messages
 - Illusion of a single system


Computing Environments – Client-Server

- Client-Server Computing
 - Dumb terminals supplanted by smart PCs
 - Many systems now **servers**, responding to requests generated by **clients**
 - **Compute-server system** provides an interface for client to request services (i.e., database)
 - **File-server system** provides an interface for clients to store and retrieve files



Computing Environments - Peer-to-Peer

- Another model of distributed system
- P2P does not distinguish clients and servers
 - Instead, all nodes are considered peers
 - May each act as client, server, or both
 - Node must join P2P network
 - Registers its service with a central lookup service on the network, or
 - Broadcast requests for service and respond to requests for service via **discovery protocol**
- Examples include Napster and Gnutella, **Voice over IP (VoIP)** such as Skype



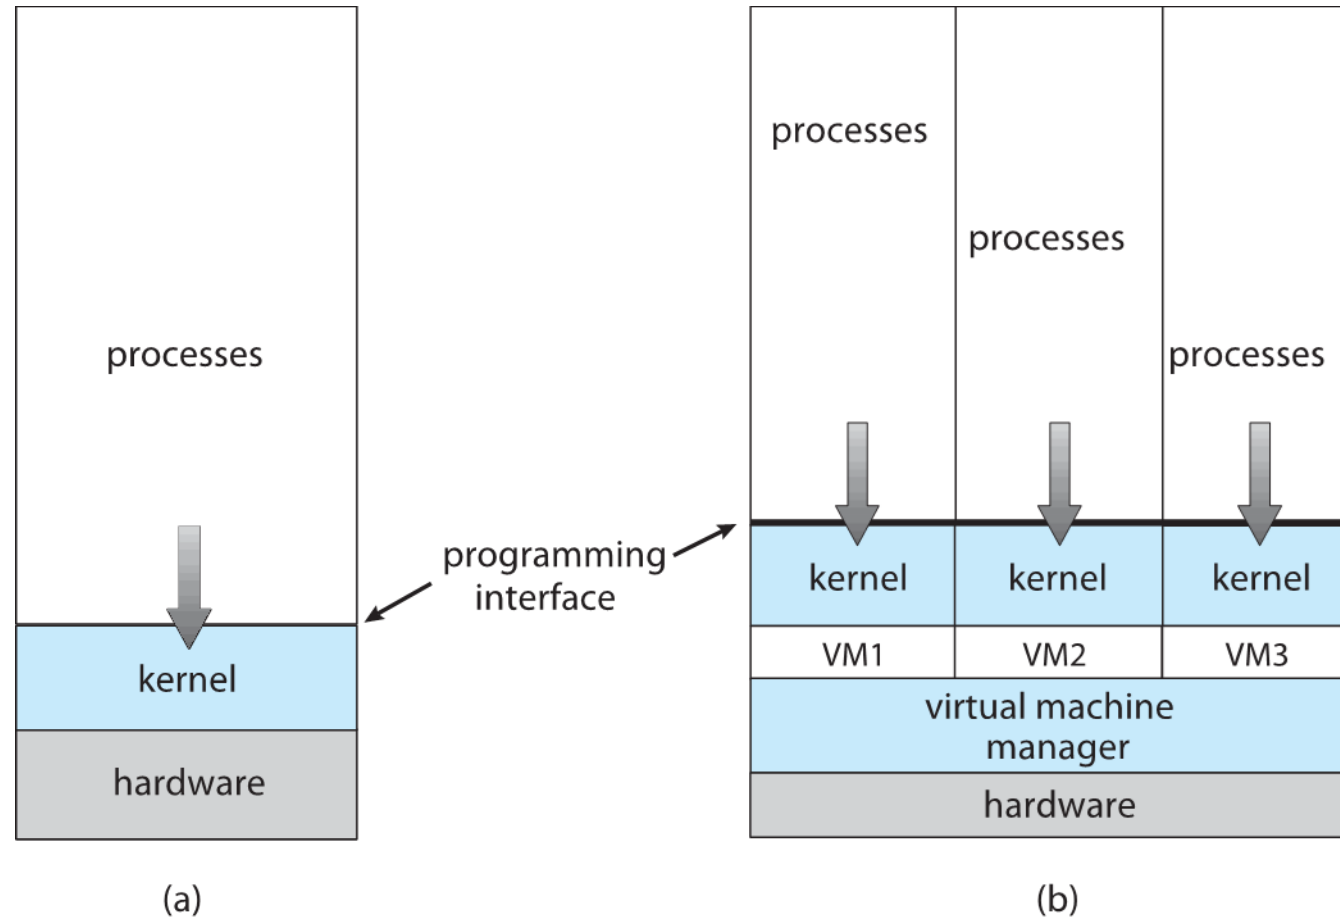
Computing Environments - Virtualization

- Allows operating systems to run applications within other OSes
 - Vast and growing industry
 - **Emulation** used when source CPU type is different from target type (i.e. PowerPC to Intel x86)
 - Generally slowest method
 - When computer language not compiled to native code – **Interpretation**
 - **Virtualization** – OS natively compiled for CPU, running **guest** OSes also natively compiled
 - Consider VMware running WinXP guests, each running applications, all on native WinXP **host** OS
 - **VMM** (virtual machine Manager) provides virtualization services
-

Computing Environments - Virtualization

- Use cases involve laptops and desktops running multiple OSes for exploration or compatibility
 - Apple laptop running Mac OS X host, Windows as a guest
 - Developing apps for multiple OSes without having multiple systems
 - QA testing applications without having multiple systems
 - Executing and managing compute environments within data centers
- VMM can run natively, in which case they are also the host
 - There is no general-purpose host then (VMware ESX and Citrix XenServer)

Computing Environments - Virtualization




Computing Environments – Cloud Computing

- Delivers computing, storage, and even apps as a service across a network
 - Logical extension of virtualization because it uses virtualization as the base for its functionality.
 - Amazon **EC2** has thousands of servers, millions of virtual machines, petabytes of storage available across the Internet, pay-based on usage
 - Many types
 - **Public cloud** – available via Internet to anyone willing to pay
 - **Private cloud** – run by a company for the company's use
 - **Hybrid cloud** – includes both public and private cloud components
 - Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e., word processor)
 - Platform as a Service (**PaaS**) – software stack ready for application use via the Internet (i.e., a database server)
 - Infrastructure as a Service (**IaaS**) – servers or storage available over the Internet (i.e., storage available for backup use)
-

Computing Environments – Cloud Computing

- Cloud computing environments composed of traditional OSES, plus VMMs, and cloud management tools
 - Internet connectivity requires security like firewalls
 - Load balancers spread traffic across multiple applications



Computing Environments – Real-Time Embedded Systems

- Real-time embedded systems most prevalent form of computers
 - Vary considerable, special purpose, limited purpose OS, **real-time OS**
 - Use expanding
 - Many other special computing environments as well
 - Some have OSes, some perform tasks without an OS
 - Real-time OS has well-defined fixed time constraints
 - Processing **must** be done within the constraint
 - Correct operation only if constraints met
-

Open-Source Operating Systems

- Operating systems made available in source-code format rather than just binary **closed-source**
- Counter to the **copy protection** and **Digital Rights Management (DRM)** movement
- Started by the **Free Software Foundation (FSF)**, which has a “copyleft” **GNU Public License (GPL)**
- Examples include **GNU/Linux** and **BSD UNIX** (including the core of **Mac OS X**), and many more
- Can use VMM like VMware Player (Free on Windows), Virtualbox (open source and free on many platforms - <http://www.virtualbox.com>)
 - Use to run guest operating systems for exploration

*Thank
you*

